

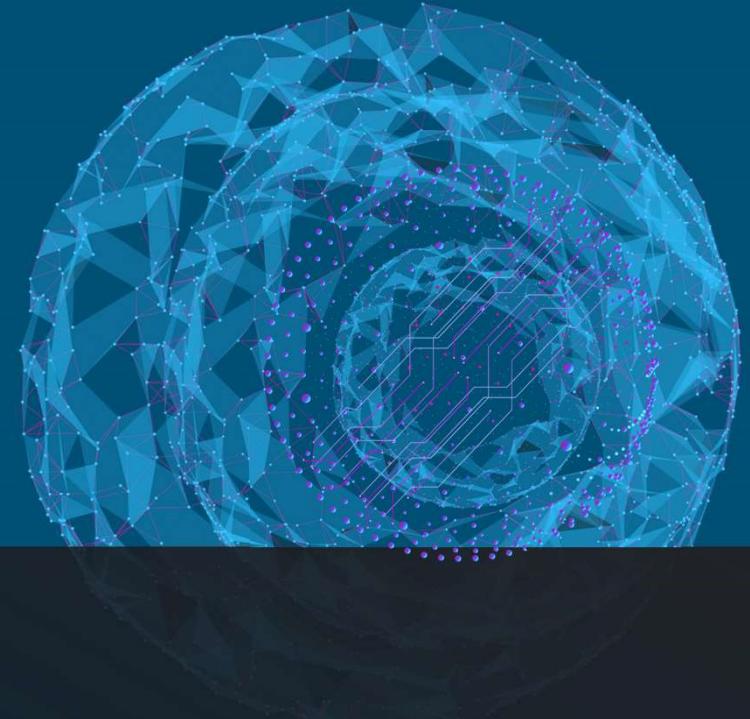
COMPUTATIONAL IMAGING

L 2. Image Processing and
Computational Illumination

Dmitry Dylov

Associate Professor

Skoltech



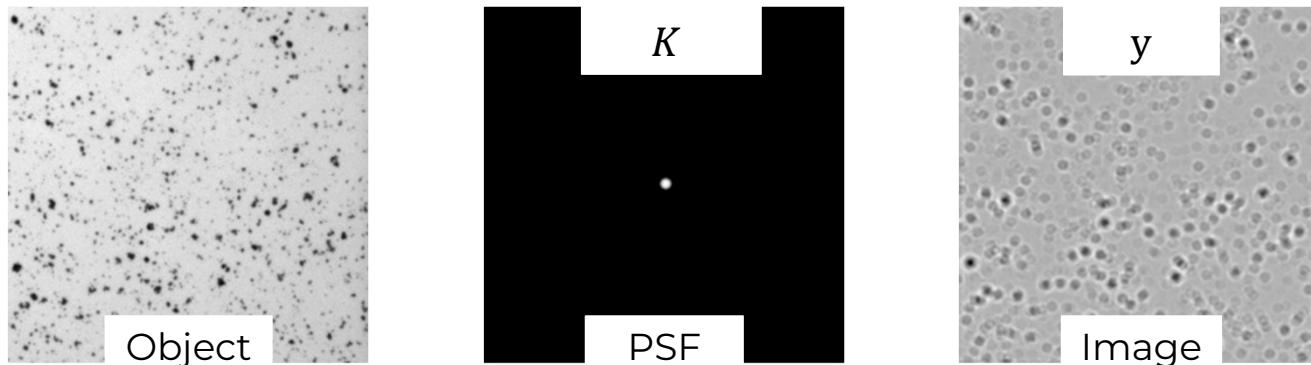
Week 1 – Understand Imaging Math & Camera Basics

Week 2 – Image Processing & Computational Illumination

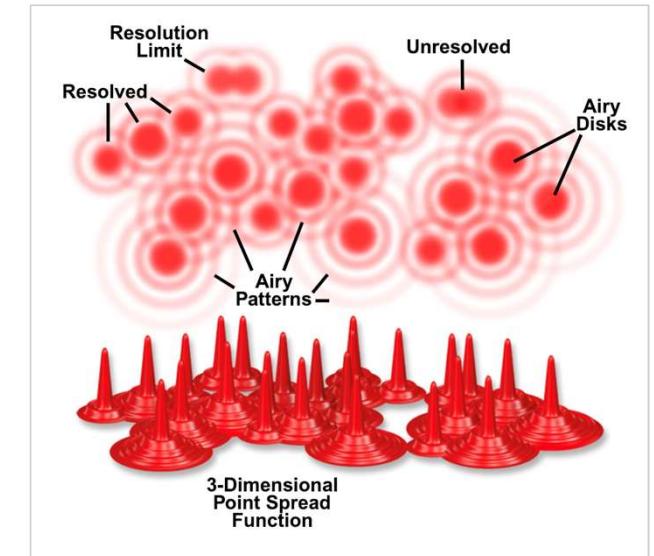
- Smoothing and Sharpening
- Filtering with convolution kernels
- Edge enhancement (Canny Filter)
- Computational Illumination:
 - structured light
 - multi-flash exposure for Canny

Week 3 – Specialized Cameras & Smart Control

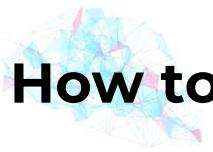
Example: Microscopy & Gaussian PSF



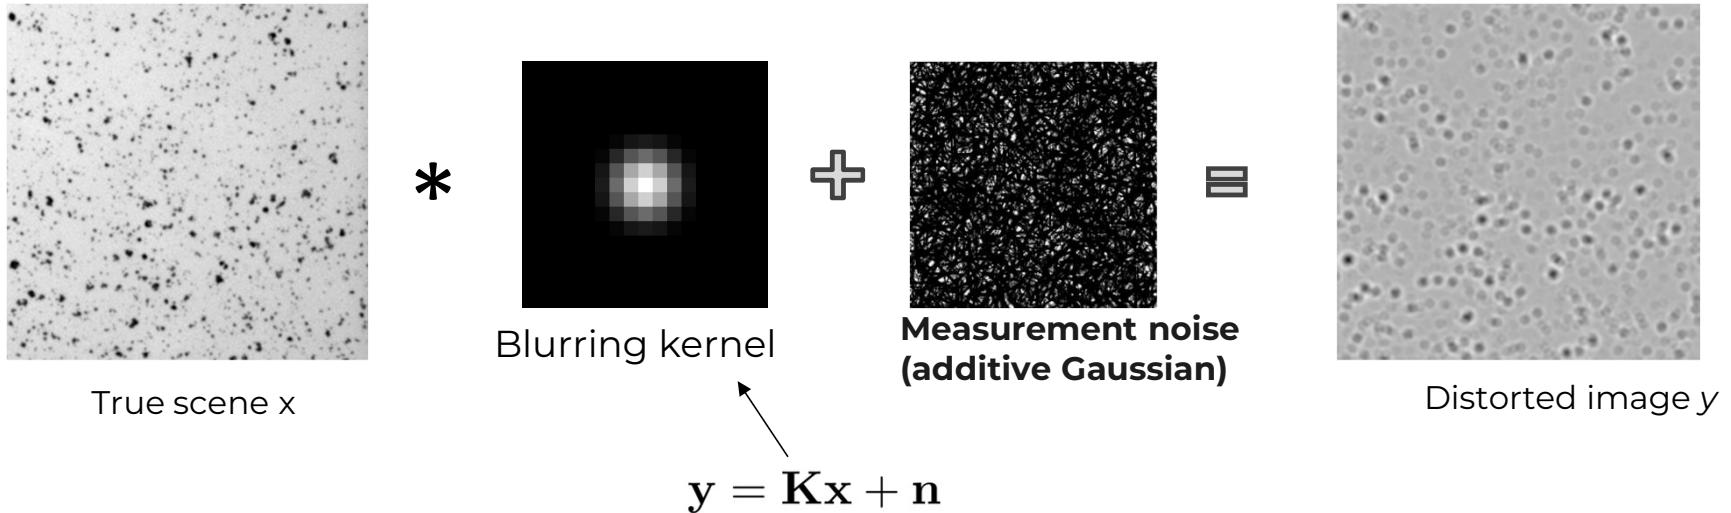
$$f(x, y) = \text{FT}^{-1} \left[\frac{Y(u, v)}{K(u, v)} \right]$$



Resolution = $0.61 \lambda/NA$
 λ = color
NA = Numerical Aperture



How to de-blur?



Solution of the optimization problem is obtained by minimizing the objective function

$$\hat{x} = \operatorname{argmin}_x \frac{1}{2} \|y - Kx\|_2^2$$

Basic Deblurring by Deconvolution

```
import numpy as np
import matplotlib.pyplot as plt

from skimage import color, data, restoration

astro = color.rgb2gray(data.astronaut())
from scipy.signal import convolve2d as conv2
psf = np.ones((5, 5)) / 25
astro = conv2(astro, psf, 'same')
astro += 0.1 * astro.std() * np.random.standard_normal(astro.shape)

deconvolved, _ = restoration.unsupervised_wiener(astro, psf)

fig, ax = plt.subplots(nrows=1, ncols=2, figsize=(8, 5),
                      sharex=True, sharey=True)

plt.gray()

ax[0].imshow(astro, vmin=deconvolved.min(), vmax=deconvolved.max())
ax[0].axis('off')
ax[0].set_title('Data')

ax[1].imshow(deconvolved)
ax[1].axis('off')
ax[1].set_title('Self tuned restoration')

fig.tight_layout()

plt.show()
```

before



after



Know blurring kernel? → Deconvolution
No? → Blind Deconvolution (iterative)

https://scikit-image.org/docs/dev/auto_examples/filters/plotrestoration.html



Another Example: Camera shake as a convolution process

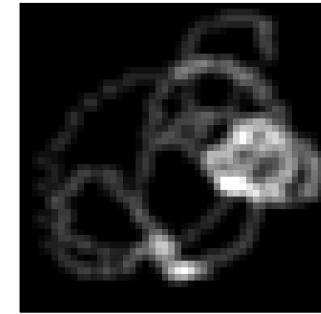
Camera shake



=



*



=



*



Bokeh experiments

Hope you will enjoy the lab!

Bokeh: Blur in out-of-focus regions of an image.



Filtering in Fourier Domain

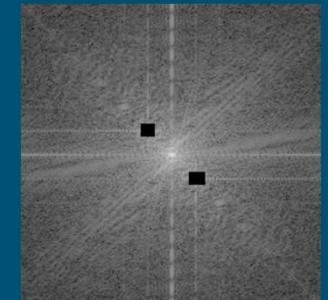
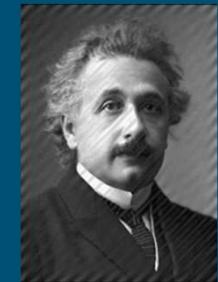
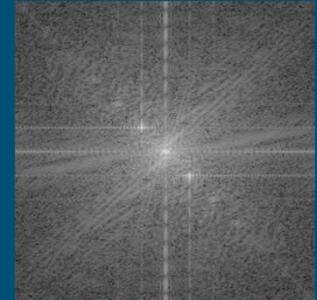
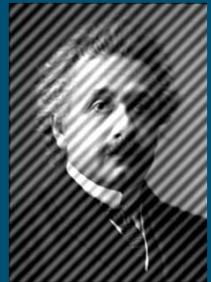
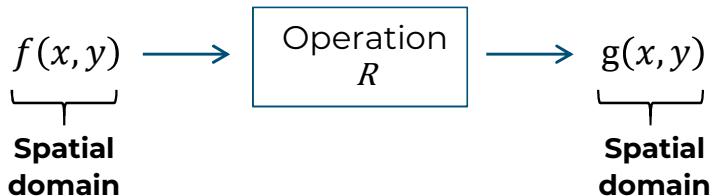


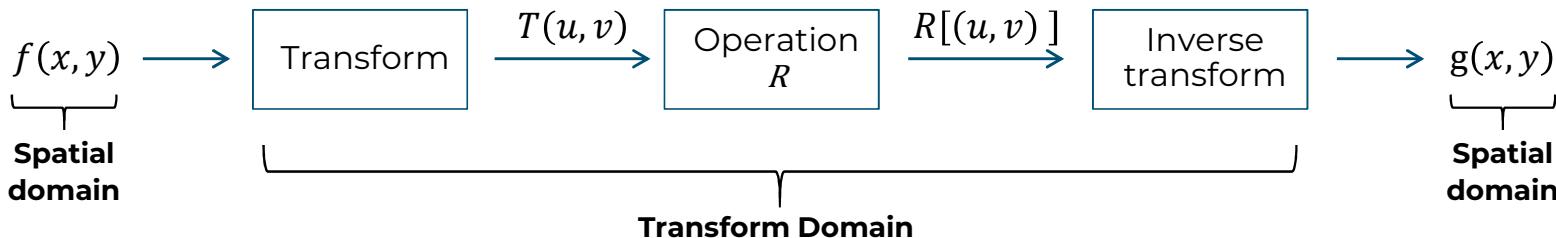


Image Filtering in the Transform Domain

- Spatial Domain



- Frequency Domain (i.e., uses Fourier Transform)

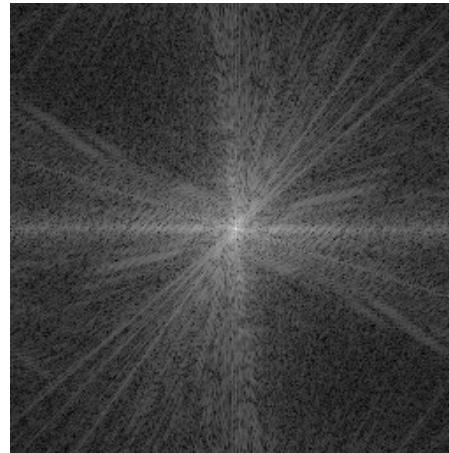




Linear filtering and convolution



$$\log(1 + |F(u, v)|)$$



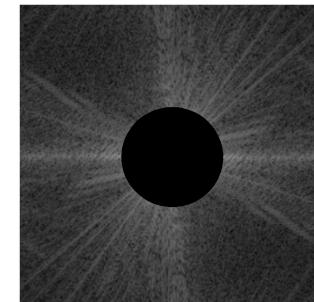


Linear filtering and convolution

- As an example, suppose $h(x, y)$ corresponds to a linear filter with frequency response defined as follows:

$$H(u, v) = 0 \text{ for } u^2 + v^2 < R^2, \quad H(u, v) = 1 \text{ otherwise}$$

- Removes low frequency components of the image
(Remember the rectangular slit calculation?)

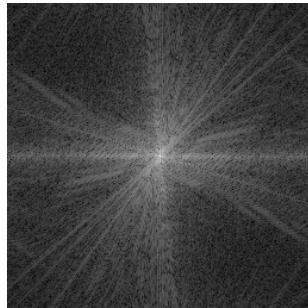




Fourier Filtering vs Convolution

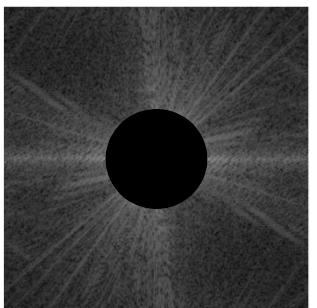


DFT

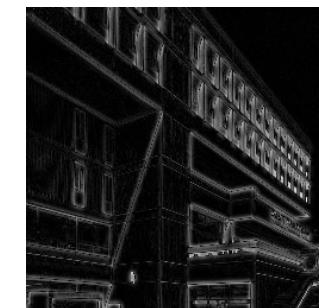
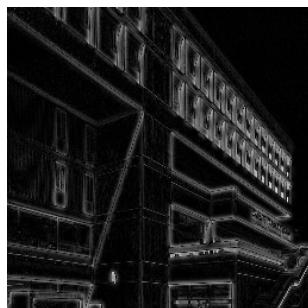


Convolve with Edge Kernel

$$H = \begin{pmatrix} 1 & 0 & -1 \\ 1 & 0 & -1 \\ 1 & 0 & -1 \end{pmatrix}$$



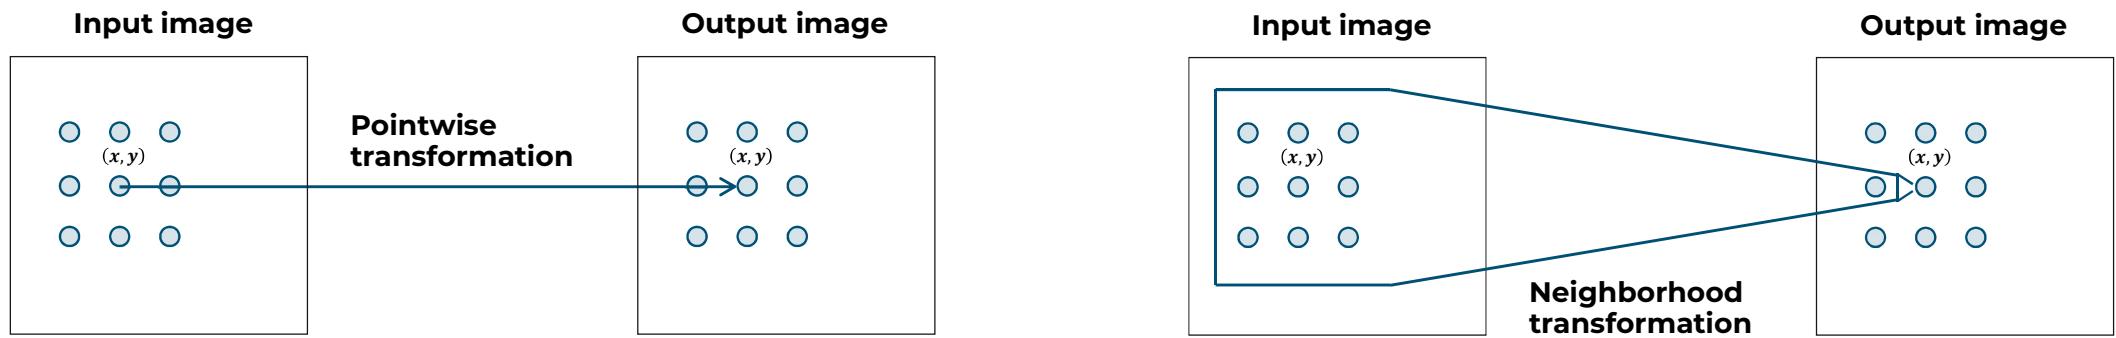
IDFT



Digital Filtering with Convolution



Image Filtering





Convolution (now, digitally)

H

$f(x, y)$

$$g(x, y) = h(x, y) * f(x, y)$$

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v]F[i - u, j - v]$$

Convolution is **commutative** and **associative**

Now that you know what it does, use convolution for various digital image processing.



Image filtering

- Modify the pixels in an image based on some function of a local neighborhood of each pixel

10	5	3
4	5	1
1	1	7

Local image data

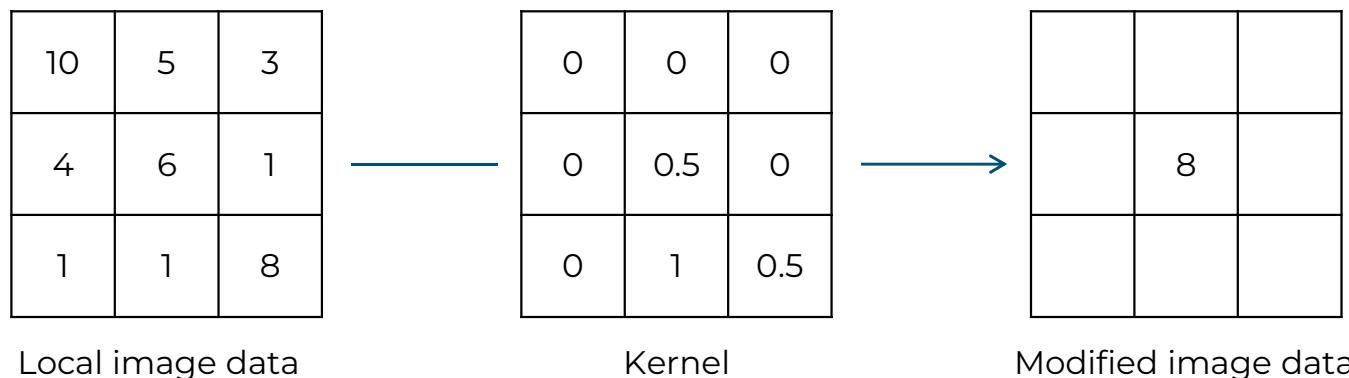
Some function

	7	

Modified image data

Linear filtering

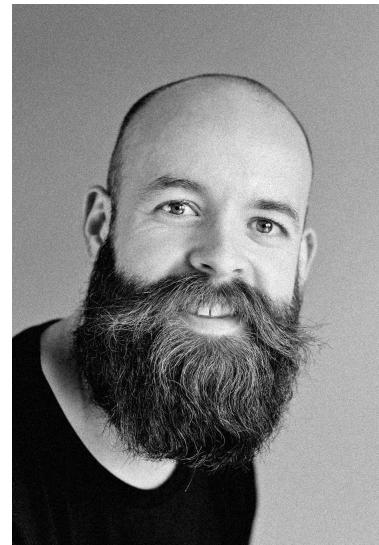
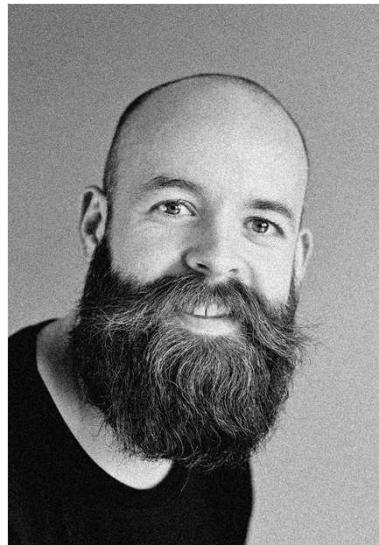
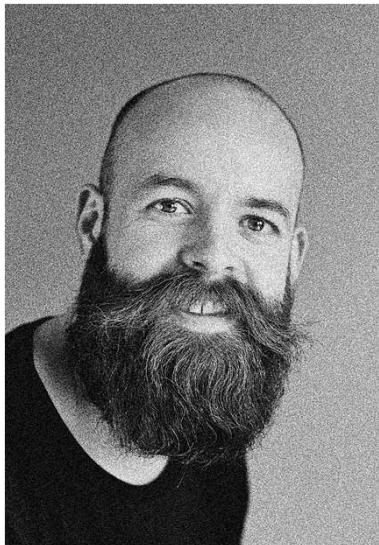
- One simple version: linear filtering (cross-correlation, convolution)
 - Replace each pixel by a linear combination of its neighbors
- The recipe for the linear combination is called the “kernel” (or “mask”, “filter”) – SAME ACTION AS PSF!





Question: Noise reduction

- Given a camera and a still scene, how can you reduce noise?



Take lots of images and average them!

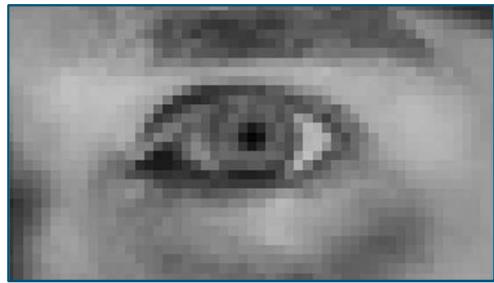
Mean filtering

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix} H$$

$$\begin{array}{c} * \\ \hline \end{array} \begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 90 & 90 & 90 & 90 & 90 & 0 & 0 \\ 0 & 0 & 0 & 90 & 90 & 90 & 90 & 90 & 0 & 0 \\ 0 & 0 & 0 & 90 & 90 & 90 & 90 & 90 & 0 & 0 \\ 0 & 0 & 0 & 90 & 90 & 90 & 90 & 90 & 0 & 0 \\ 0 & 0 & 0 & 90 & 0 & 90 & 90 & 90 & 0 & 0 \\ 0 & 0 & 0 & 90 & 90 & 90 & 90 & 90 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 90 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix} = \begin{bmatrix} & & & & & & & & \\ 0 & 10 & 20 & 30 & 30 & 30 & 20 & 10 & \\ 0 & 20 & 40 & 60 & 60 & 60 & 40 & 20 & \\ 0 & 30 & 60 & 90 & 90 & 90 & 60 & 30 & \\ 0 & 30 & 50 & 80 & 80 & 90 & 60 & 30 & \\ 0 & 30 & 50 & 80 & 80 & 90 & 60 & 30 & \\ 0 & 20 & 30 & 50 & 50 & 60 & 40 & 20 & \\ 10 & 20 & 30 & 30 & 30 & 30 & 20 & 10 & \\ 10 & 10 & 10 & 0 & 0 & 0 & 0 & 0 & \\ & & & & & & & & \end{bmatrix} G$$



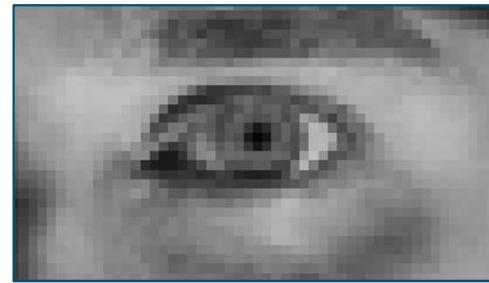
Linear filters: examples



*

0	0	0
0	1	0
0	0	0

=

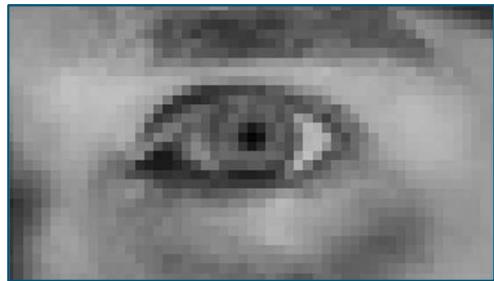


Original

Identical image



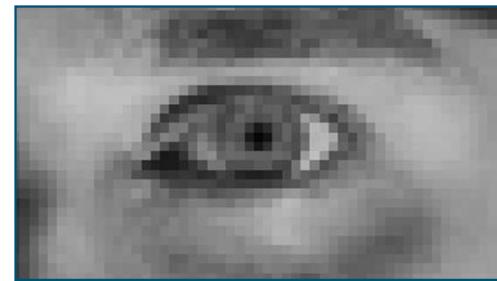
Linear filters: examples



*

0	0	0
1	0	0
0	0	0

=

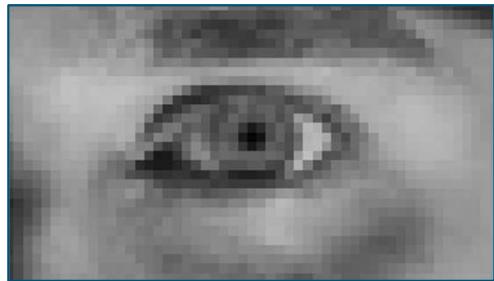


Original

Shifted left By 1 pixel



Linear filters: examples



Original

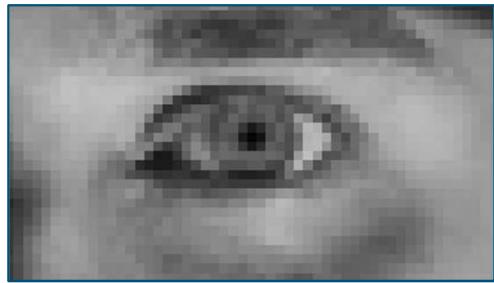
$$\ast \frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array} =$$



Blur (with a mean filter)

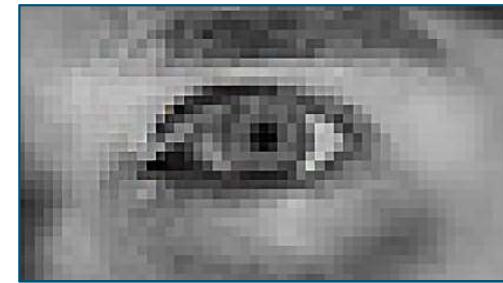


Linear filters: examples



Original

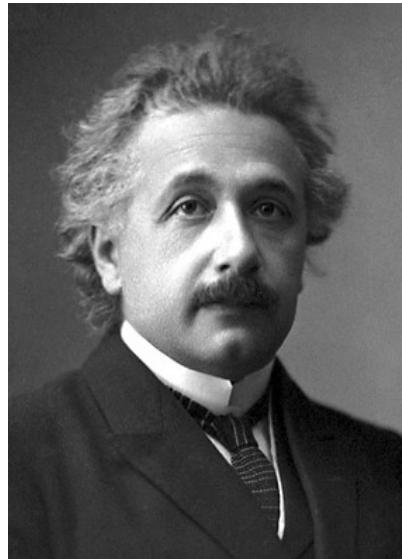
$$* \left(\begin{array}{ccc} 0 & 0 & 0 \\ 0 & 2 & 0 \\ 0 & 0 & 0 \end{array} - \frac{1}{9} \begin{array}{ccc} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{array} \right) =$$



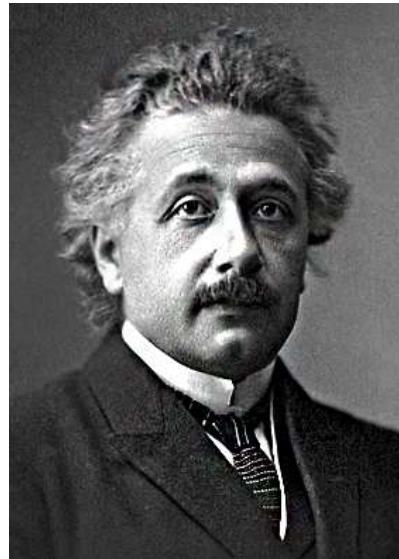
Sharpening filter
(accentuates edges)



Sharpening



Before



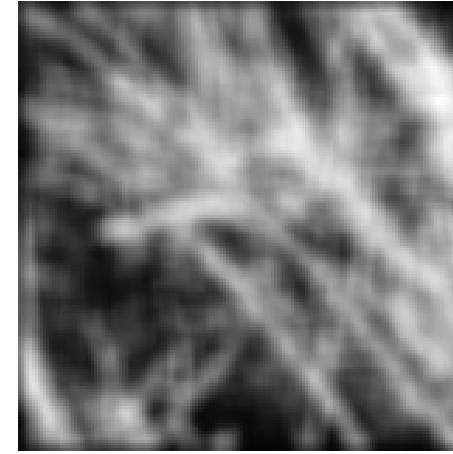
After



Smoothing with box filter revisited

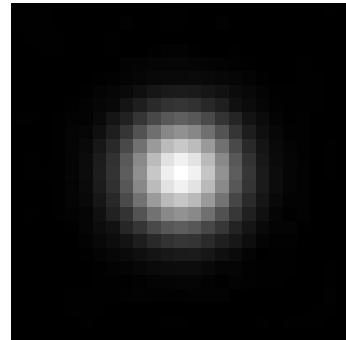
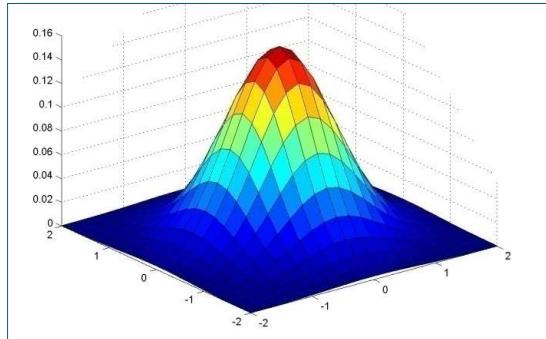


$$* \quad \boxed{\square} \quad =$$





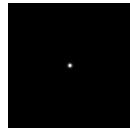
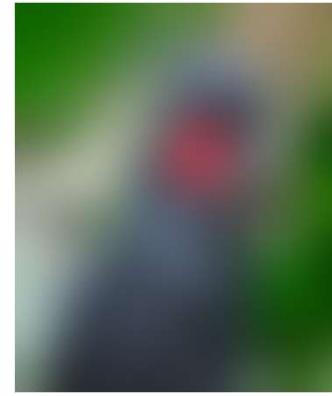
Gaussian Kernel



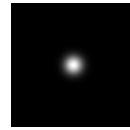
$$G_{\sigma} = \frac{1}{2\pi\sigma^2} e^{-\frac{(x^2+y^2)}{2\sigma^2}}$$



Gaussian filters



$\sigma = 1$ pixel



$\sigma = 5$ pixel



$\sigma = 10$ pixel



$\sigma = 30$ pixel

- Removes “high-frequency” components from the image (low-pass filter)
- Convolution with self is another Gaussian



Gaussian filters – optimal amount



Original



Noisy



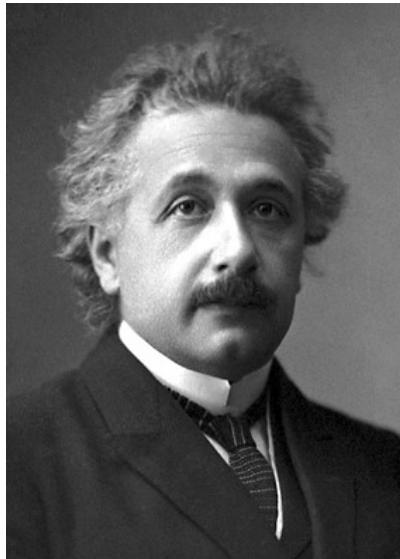
Filtered
 $\sigma = 1.5$



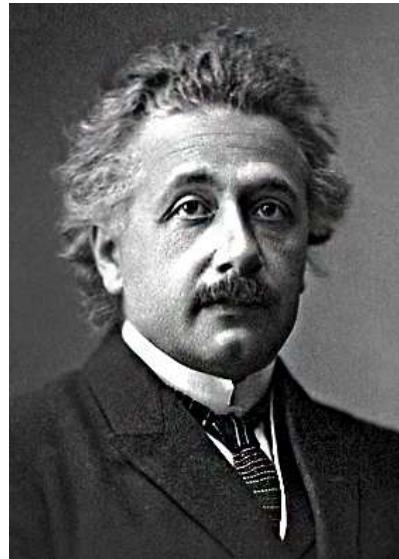
Filtered
 $\sigma = 3.0$



Sharpening



Before



After



Sharpening revisited



- What does blurring take away?



- Let's add it back

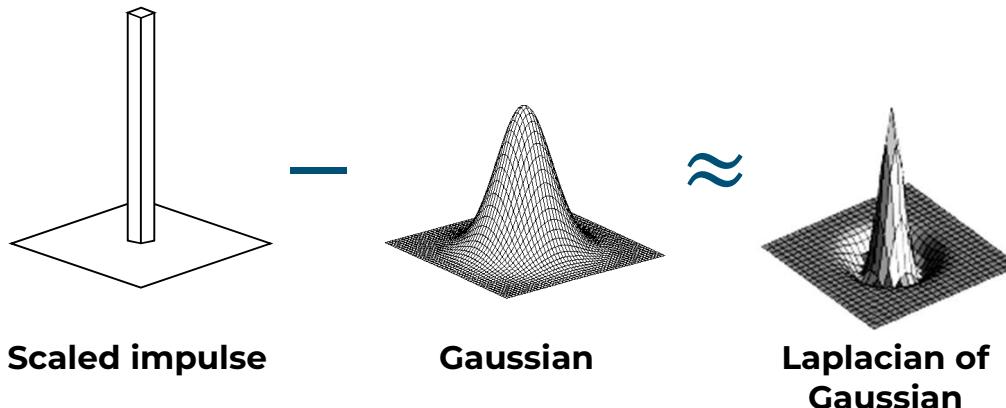


Sharpen filter

$$F + a(F - F * H) = (1 + a) - a(F * H) = F * ([1 + a]e - H)$$

↑
Image Blurred image

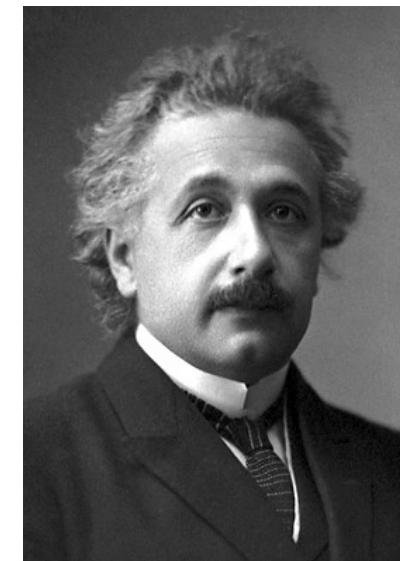
↑
Image



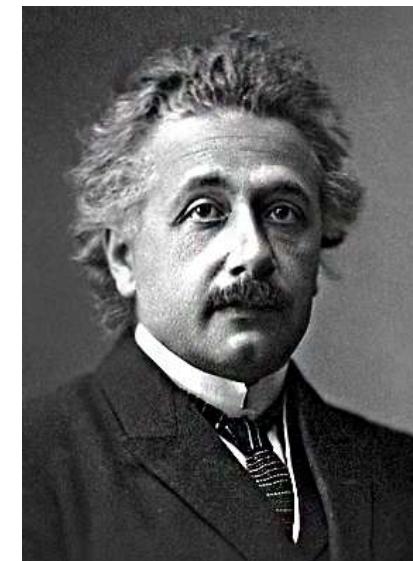
Scaled impulse

Gaussian

Laplacian of
Gaussian



Before



After

Cross-correlation



Cross-correlation

- Let F be the image, H be the kernel (of size $2k+1 \times 2k+1$), and G be the output image

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v]F[i + u, j + v]$$

This is called a **cross-correlation** operation:

$$G = H \otimes F$$

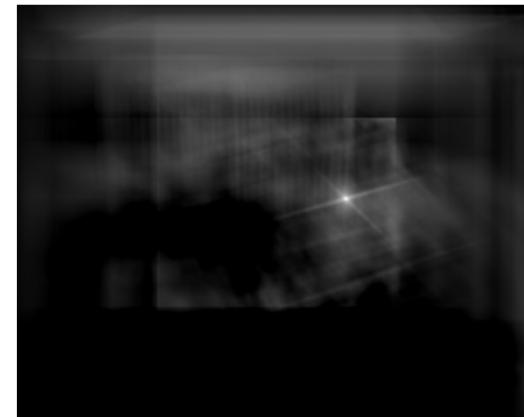


Normalized Correlation

- Measure the similarity between images or parts of images.

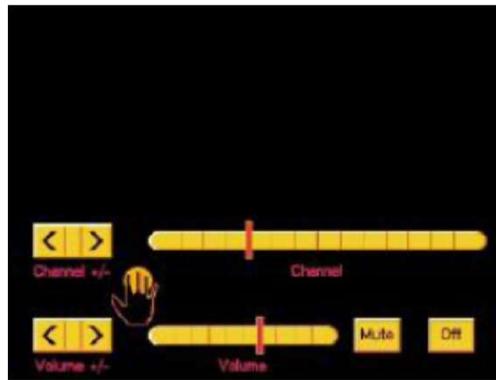


$$\otimes \quad \text{mask} =$$





Application: TV Remote Control



Template (left), image (middle), normalized correlation (right)

Note peak value at the true position of the hand

<https://www.spiedigitallibrary.org/conference-proceedings-of-spie/10452/104520Y/>



Summary: Digital Filtering

Convolution = most basic go-to linear filtering approach

Correlation = shows similarity between images

Fourier Domain = sometimes it is easier to know what to Filter there



Image Enhancement 2: Edge Detection

Sharpen filter

$$F + a(F - F * H) = F(1 + a) - a(F * H) = F * ([1 + a]e - H)$$

↑
Image Blurred image

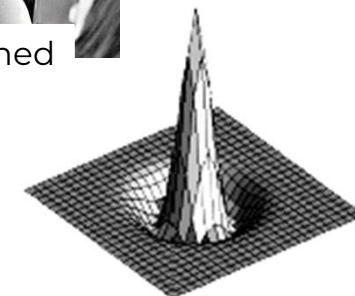
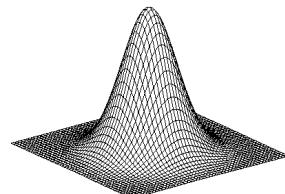


+ α



=

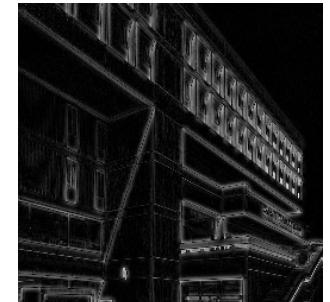
Unit impulse
(identity)



Convolve with Edge Kernel

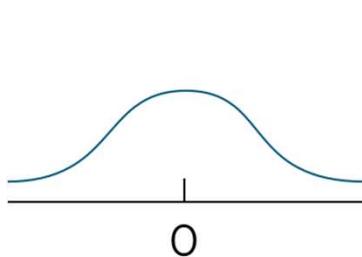
(high-pass filter)

$$H = \begin{pmatrix} 1 & 0 & -1 \\ 1 & 0 & -1 \\ 1 & 0 & -1 \end{pmatrix}$$

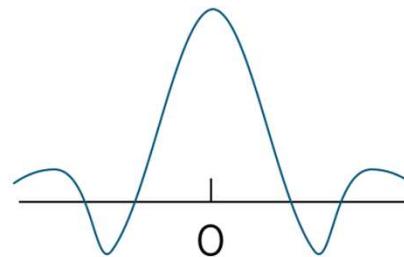
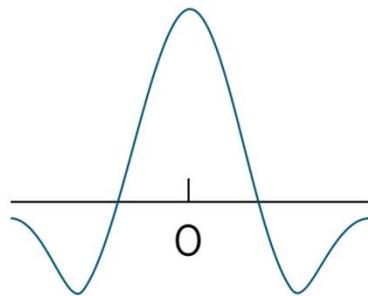


Kernel collection

- Depends on the application.
- Usually by sampling certain functions and their derivatives.



Good for image
smoothing

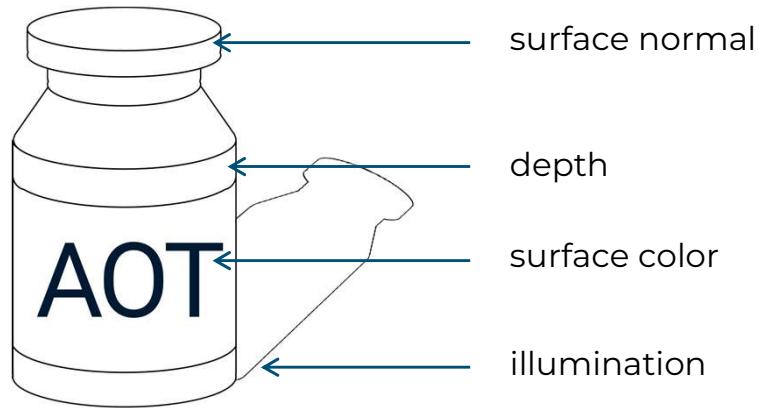


Good for image
sharpening

Edge detection



- Convert a 2D image into a set of curves
 - Extracts salient features of the scene
 - More compact than pixels



- Origins: discontinuity in...

Characterizing edges

- An edge is a place of rapid change in the image intensity function

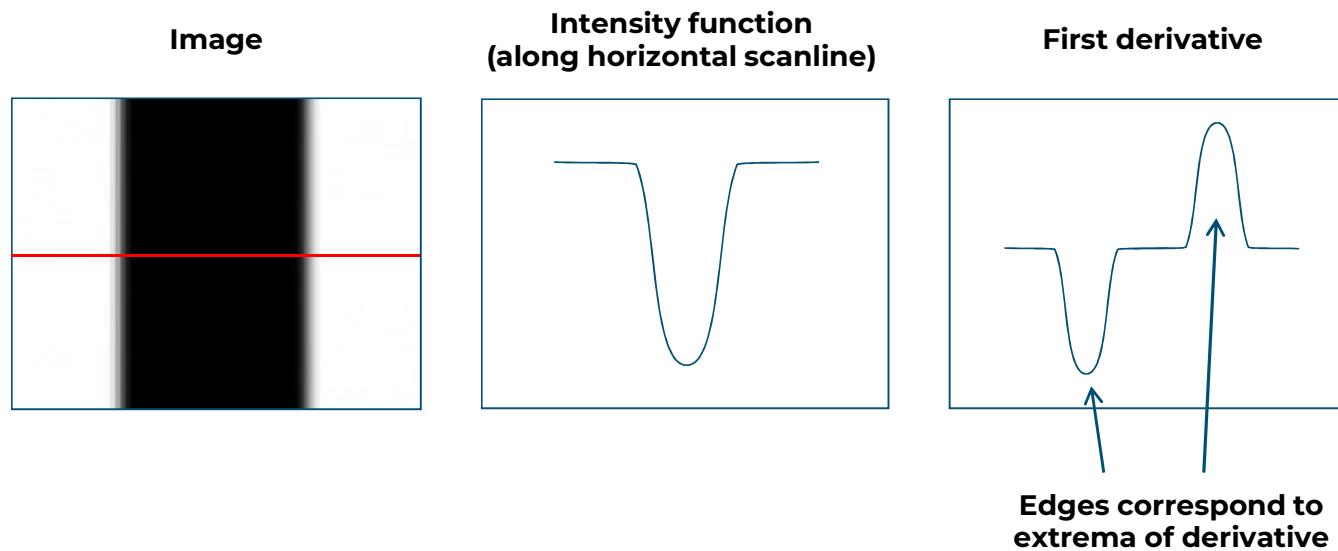


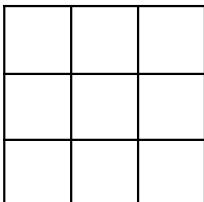
Image derivatives

- **How can we differentiate a digital image $F[x,y]$?**

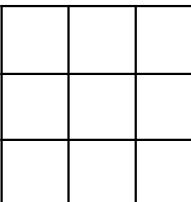
- Option 1: reconstruct a continuous image, f , then compute the derivative
- Option 2: take discrete derivative (finite difference)

$$\frac{\partial f}{\partial x}[x,y] \approx F[x+1,y] - F[x,y]$$

How would you implement this as a linear filter?

$$\frac{\partial f}{\partial x}:$$


H_x

$$\frac{\partial f}{\partial y}:$$


H_y

Image gradient

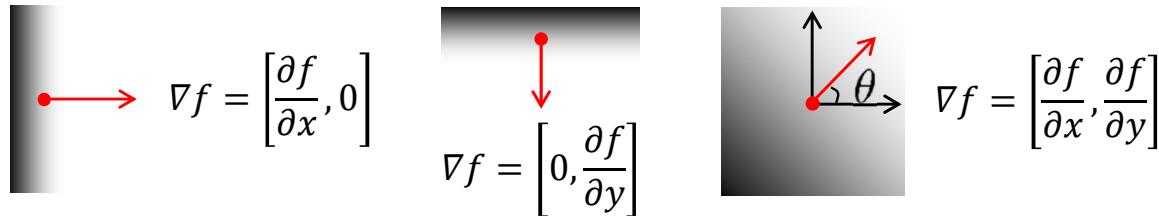
- **The gradient of an image:**

$$\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$

The edge strength is given by the gradient magnitude:

$$\|\nabla f\| = \sqrt{\left(\frac{\partial f}{\partial x}\right)^2 + \left(\frac{\partial f}{\partial y}\right)^2}$$

The gradient points in the direction of most rapid increase in intensity

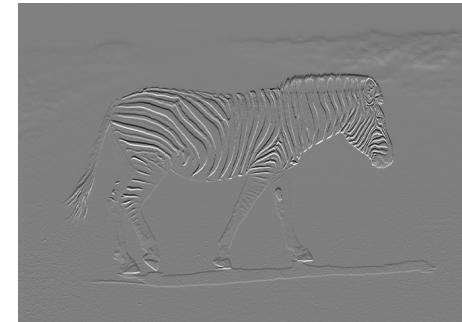
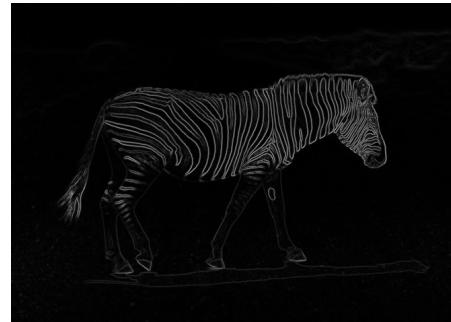
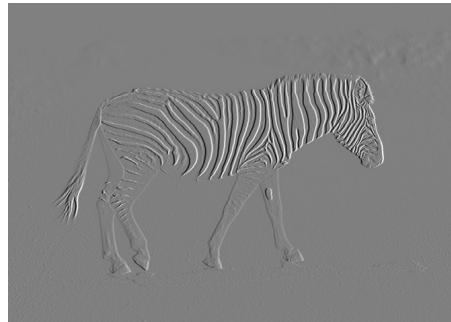


The gradient direction is given by:

$$\phi = \tan^{-1} \left(\frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$$

- **how does this relate to the direction of the edge?**

Image gradient



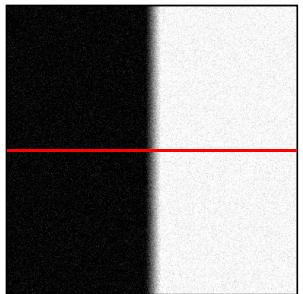
f

$$\frac{\partial f}{\partial x}$$

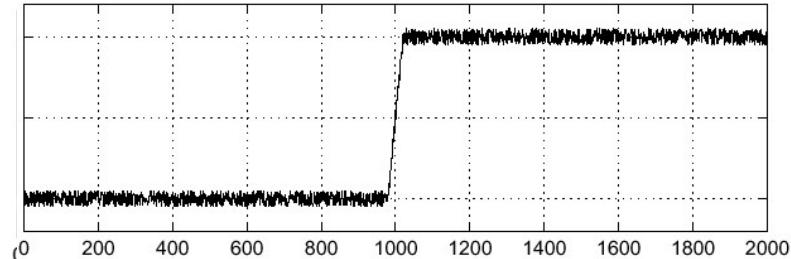
$$\|\nabla f\| = \sqrt{\left(\frac{\partial f}{\partial x}\right)^2 + \left(\frac{\partial f}{\partial y}\right)^2}$$

$$\frac{\partial f}{\partial y}$$

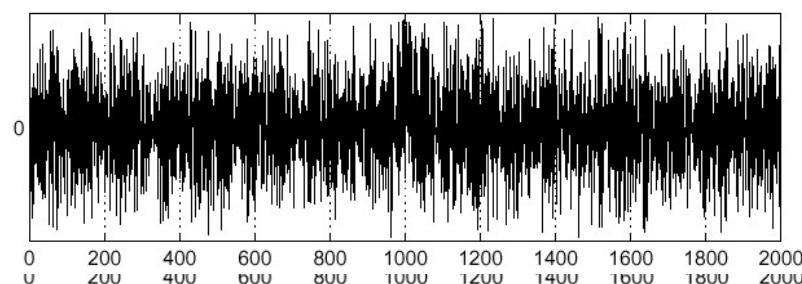
Effects of noise



Noisy input
image



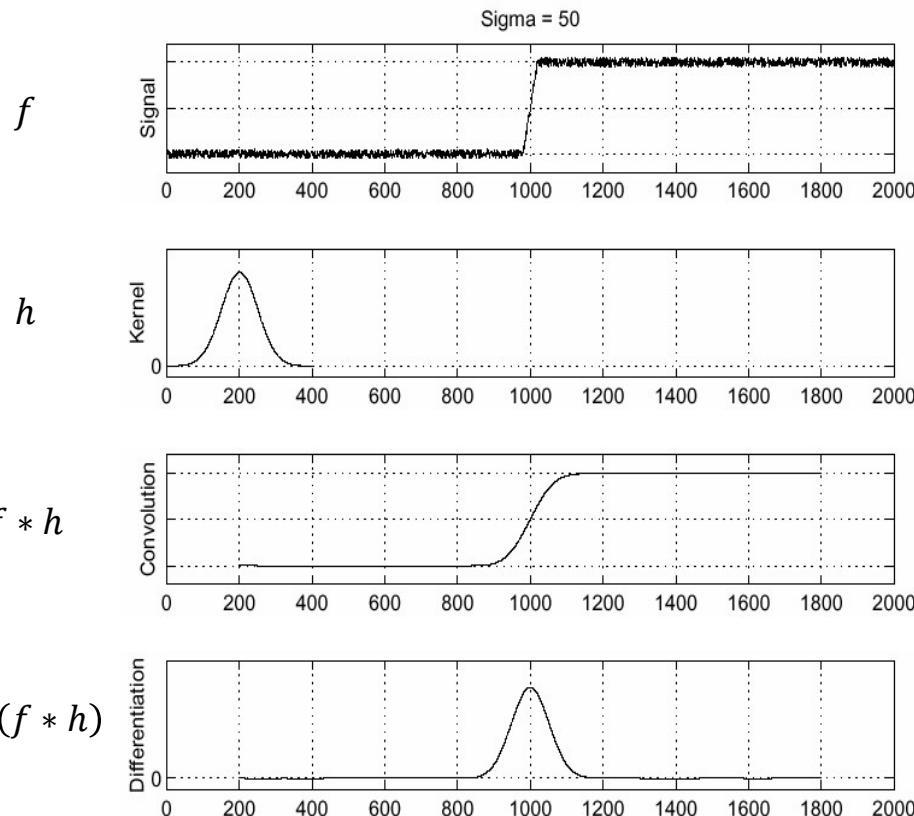
$f(x)$



$\frac{d}{dx}f(x)$

Where is the edge?

Solution: smooth first



To find edges, look for peaks in

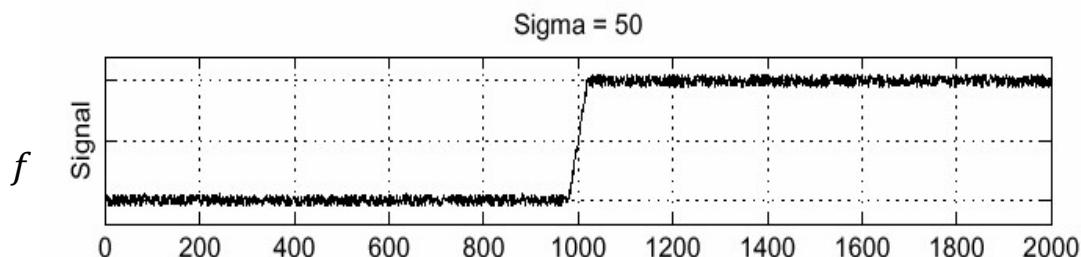
$$\frac{d}{dx}(f * h)$$

Associative property of convolution

Differentiation is convolution, and convolution is associative:

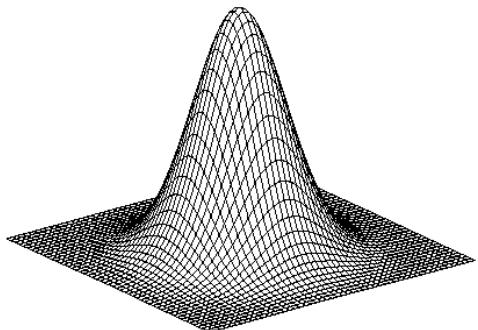
$$\frac{\partial}{\partial x} (f * h) = f * \frac{d}{dx} h$$

This saves us one operation:



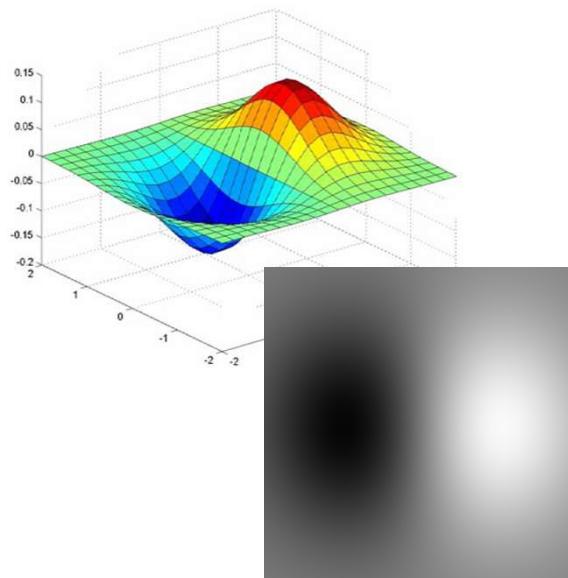
Instead, compute derivatives of a Gaussian kernel and then convolve

2D edge detection filters



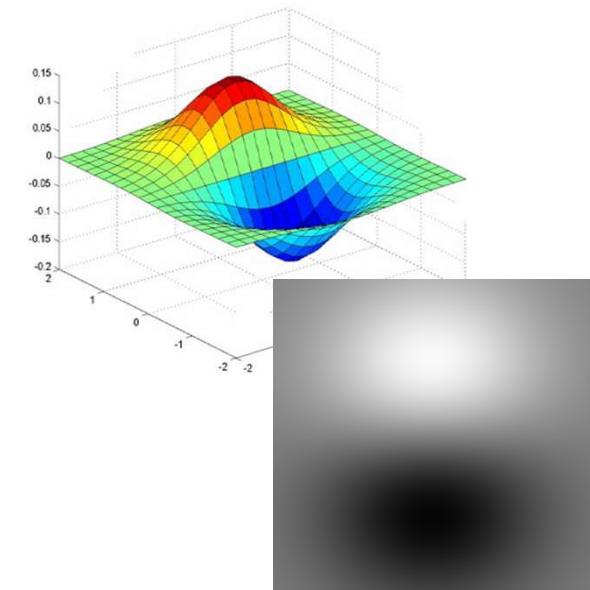
Gaussian

$$h_\sigma(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{2\sigma^2}}$$



$$\frac{\partial}{\partial x} h_\sigma(u, v)$$

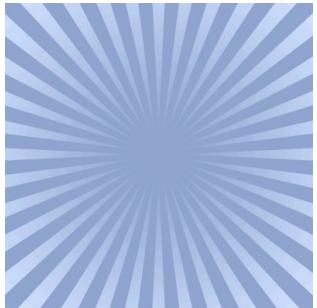
Derivative of Gaussian (x)
x-direction



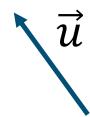
$$\frac{\partial}{\partial y} h_\sigma(u, v)$$

Derivative of Gaussian (y)
y-direction

Side note: How would you compute a directional derivative?

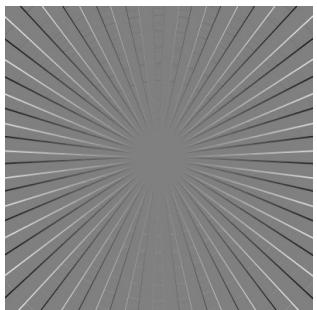


f

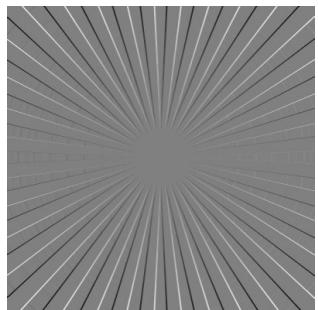


$$\nabla_{\vec{u}} f = ?$$

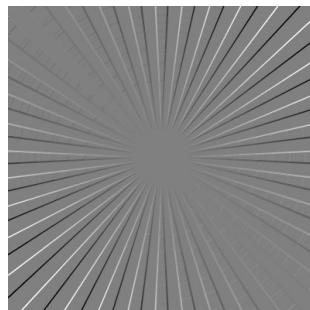
Directional derivative is a linear combination of partial derivatives



+



=

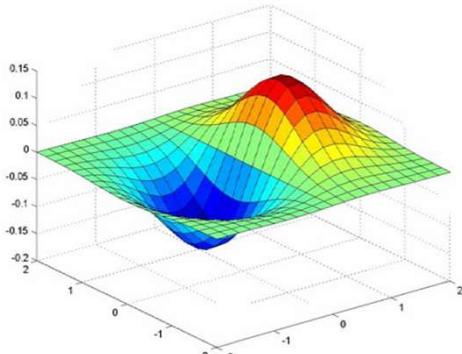


$$\frac{\partial f}{\partial x} \cdot u_x$$

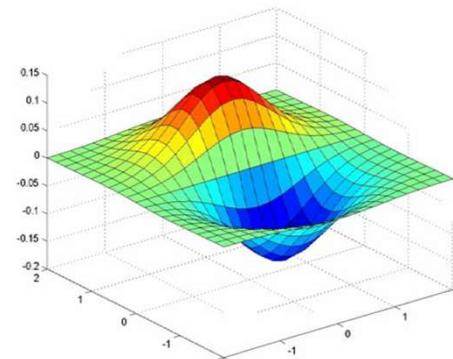
$$\frac{\partial f}{\partial y} \cdot u_y$$

$$\nabla_{\vec{u}} f$$

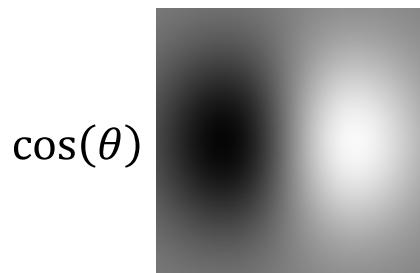
Derivative of Gaussian filter



x-direction

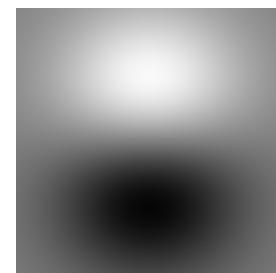


y-direction

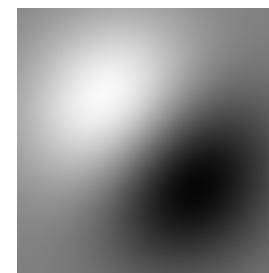


$\cos(\theta)$

$+\sin(\theta)$



=



The Sobel operator

- Common approximation of derivative of Gaussian

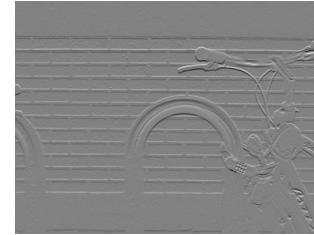
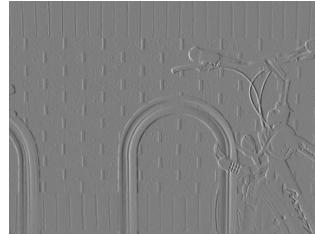
$$\frac{1}{8} \begin{matrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{matrix}$$

S_x

$$\frac{1}{8} \begin{matrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{matrix}$$

S_y

Sobel operator: example



Example



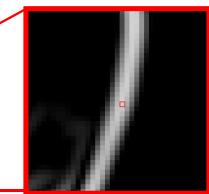
Original image (Lena Söderberg)

Finding edges



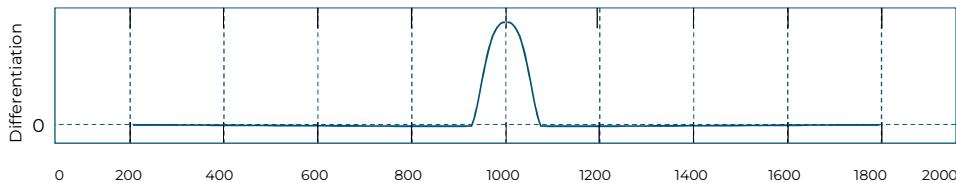
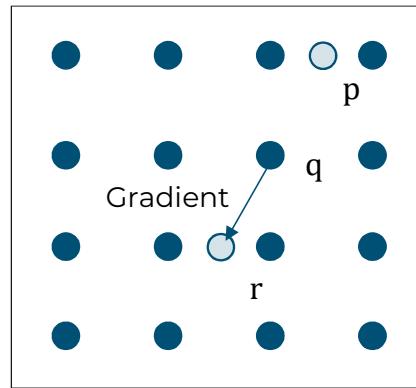
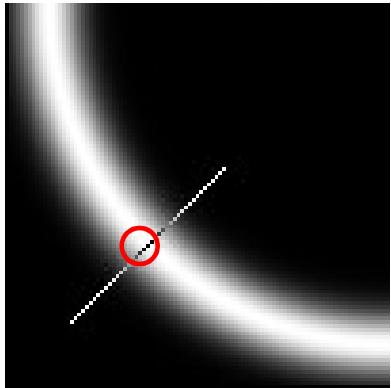
Gradient magnitude

Finding edges



thresholding

Non-maximum suppression



- Check if pixel is local maximum along gradient direction
 - requires *interpolating* pixels p and r

Finding edges



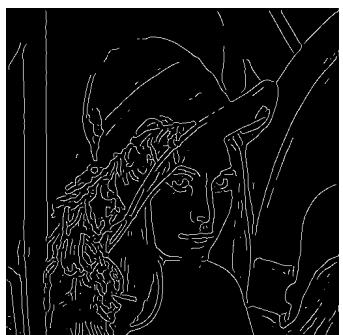
thresholding

Finding edges



thinning (non-maximum suppression)

Canny edge detector



1. Filter image with derivative of Gaussian
2. Find magnitude and orientation of gradient
3. Non-maximum suppression
4. Linking and thresholding (hysteresis):
 - Define two thresholds: low and high
 - Use the high threshold to start edge curves and the low threshold to continue them

`cv2.Canny()`

Canny edge detector



Original



Canny with $\sigma = 1$



Canny with $\sigma = 2$

- The choice of σ depends on desired behavior
 - large σ detects “large-scale” edges
 - small σ detects fine edges



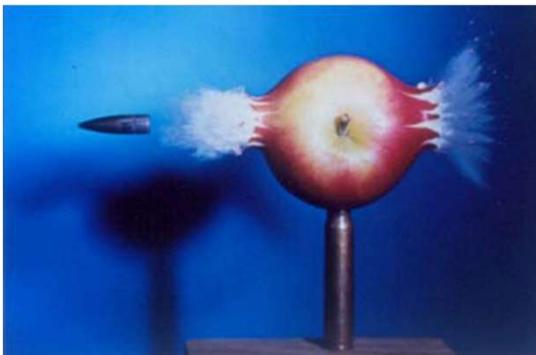
Computational Illumination

Computational Illumination

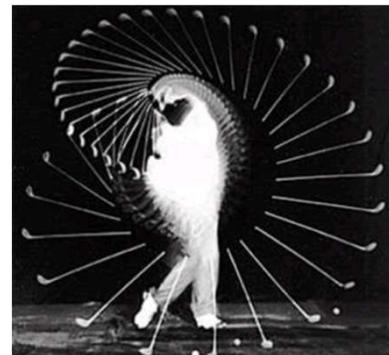
Programmable 4D Illumination Field + Time + Wavelength

- **ON or OFF, Duration, Brightness**
 - Flash/No-flash
- **Light position**
 - Relighting: Programmable dome
 - Shape enhancement: Multi-flash for depth edges
- **Light color/wavelength**
- **Spatial Modulation**
 - Synthetic Aperture Illumination
- **Temporal Modulation**
 - Motion Tracking
- **Exploiting natural lighting condition**
 - Day/Night Fusion, Time Lapse, Glare

Edgerton, 1930 (MIT museum)



**Stroboscope
(Electronic Flash)**



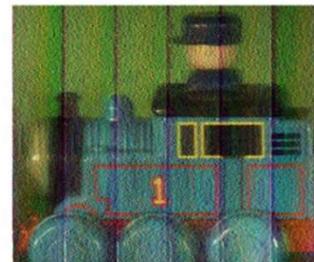
**Multi-Flash Sequential
Photography**





Flutter Shutter: deconvolution in time

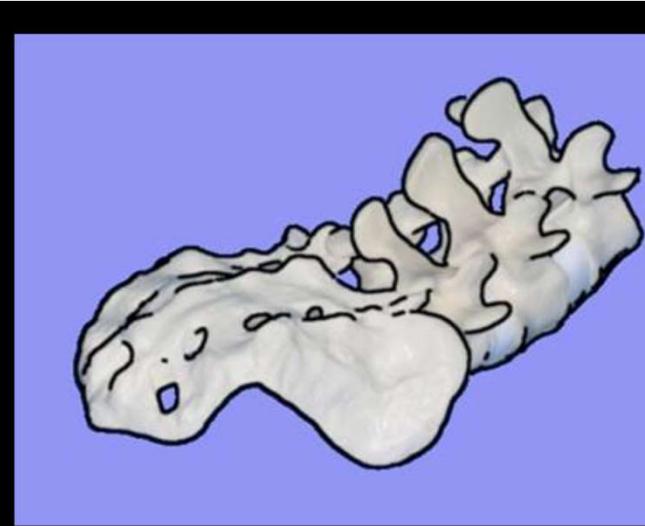
Long Exposure



Coded sequence





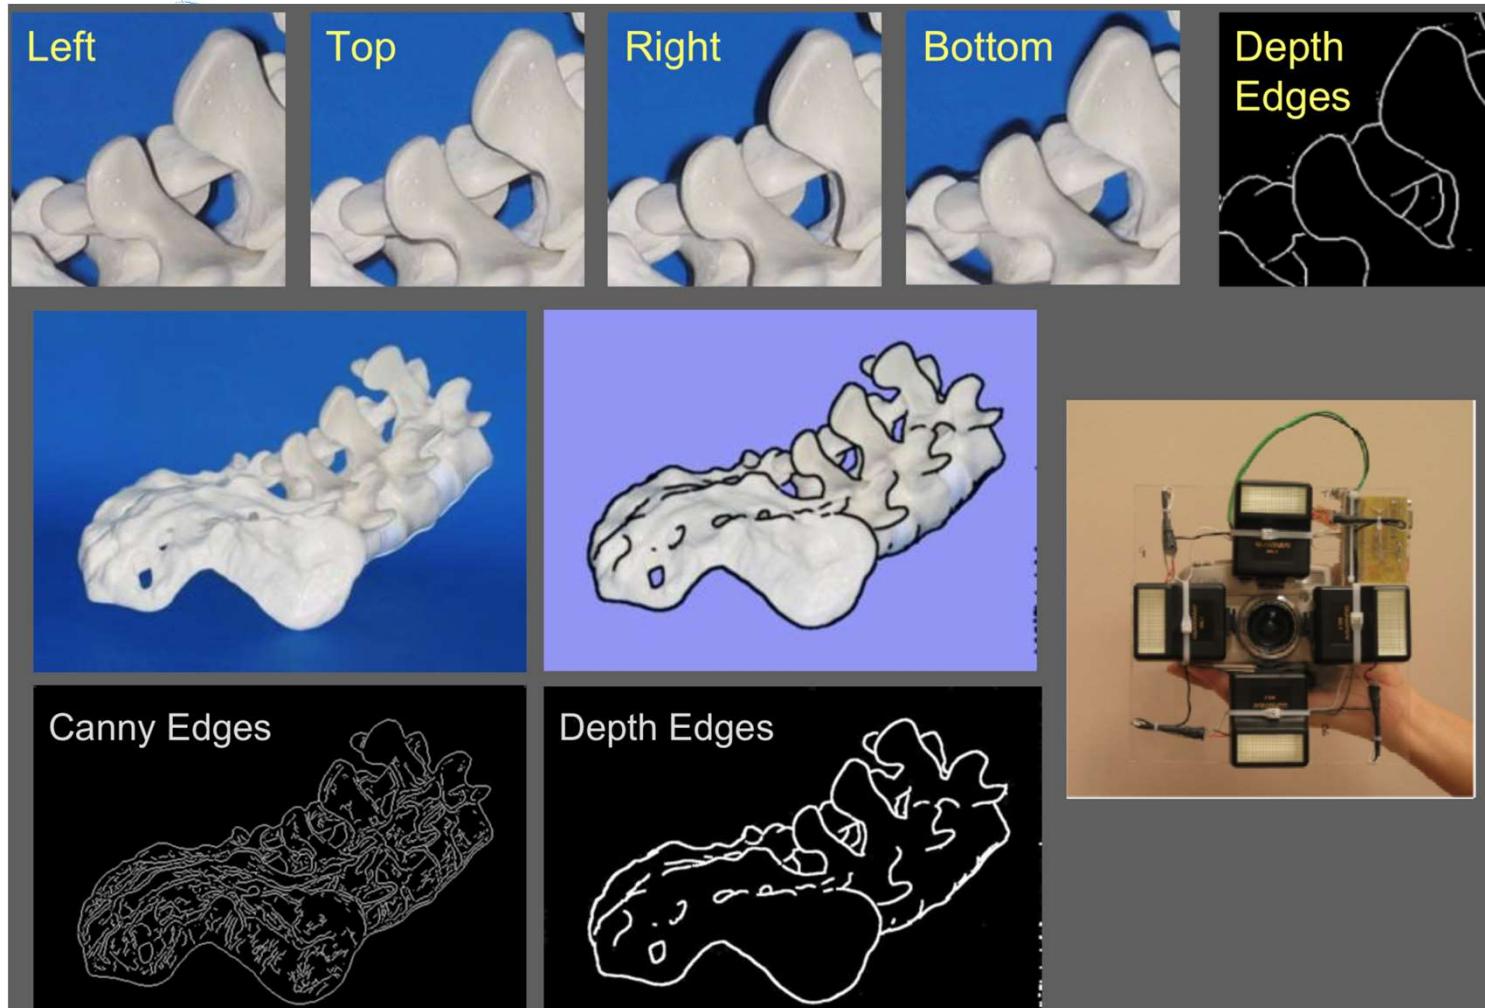


Canny

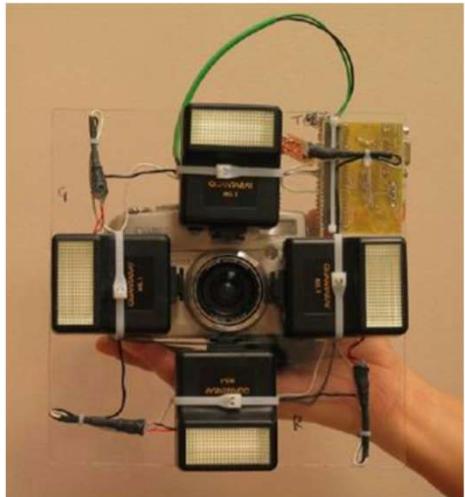


Our Method





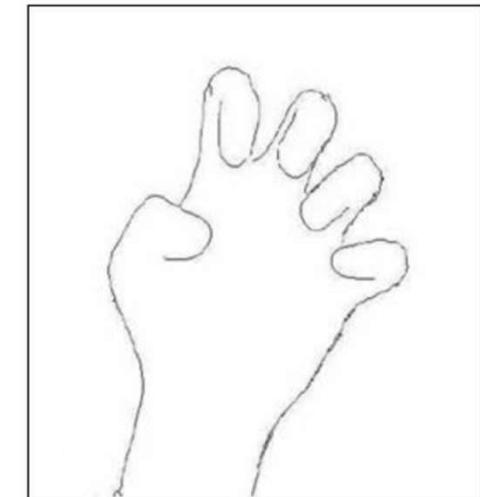
Example of computational illumination: Illuminate 4 times, then reconstruct



Input Photo



Canny Edges



Milti-flash camera image



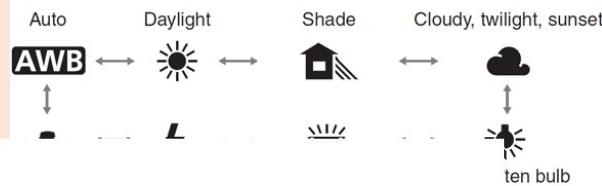
Color Imaging

Color imaging



“The Dress”

White Balance



Untitled Poll

00:01:35 | 1 question | 21 of 27 (77%) participated

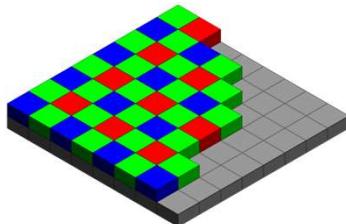
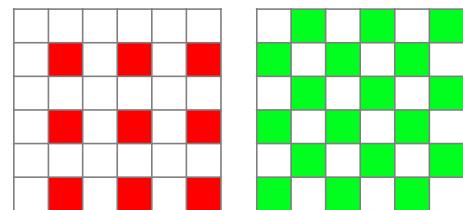
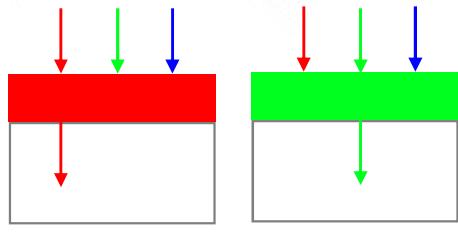
1. Which color is the dress (Single Choice) *

21/21 (100%) answered

Gold (9/21) 43%

Blue (12/21) 57%

Color sensor



Bayer Filter

Incoming light

Filter layer

Sensor array

Resulting pattern

De-mosaicking



Incoming light

Filter layer

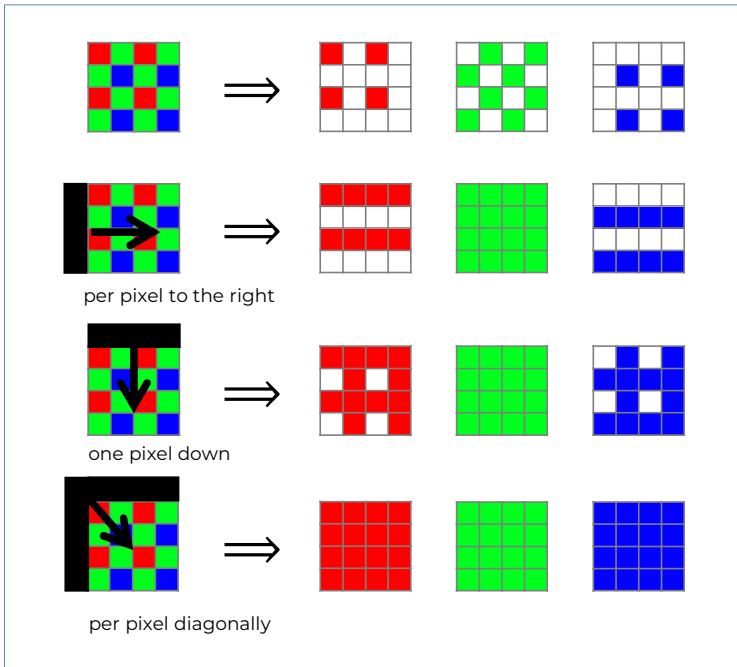
Sensor array

Resulting pattern

Interpolate missing colors!
Noise? -- Average

Pixel Shifting: in every mobile phone today

Pixel shifting



**(your natural handshake
is useful for averaging)**

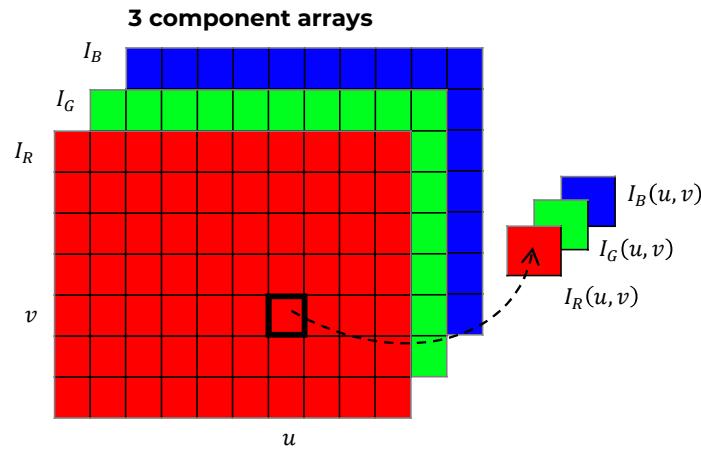
Color Images

- True color: Uses all colors in color space
- Indexed color: Uses only some colors
- Which subset of colors to use? Depends on application
- True color:
 - used in applications that contain many colors with subtle differences
 - e.g., digital photography or photorealistic rendering
 - Component ordering
 - Packed ordering

True Color: Ordering

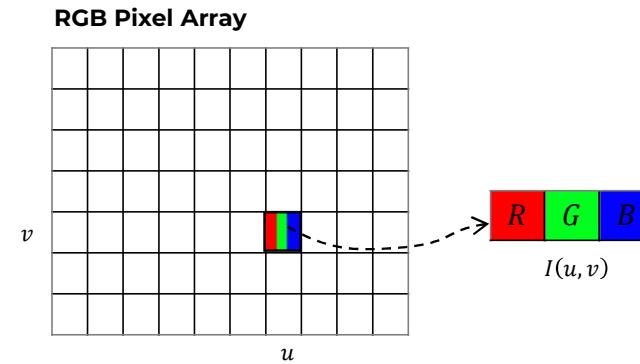
Component Ordering

Colors in 3 separate arrays of similar length
Retrieve same location (u, v) in each R, G and B array



Packed Ordering

Component (RGB) values representing a pixel's color is packed together into single element



Indexed Images

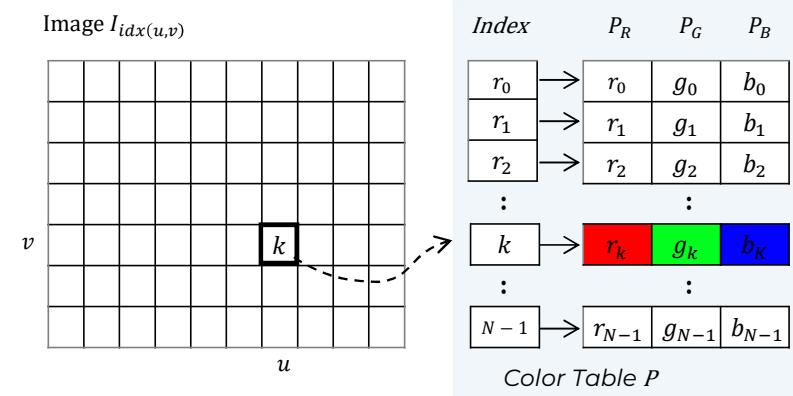
- Permit only limited number of distinct colors ($N = 2$ to 256)
- Used in illustrations or graphics containing large regions with same color
- Instead of intensity values, image contains indices into color table or palette
- Palette saved as part of image
- Converting from true color to indexed color requires quantization

RGB full-color images (24-bit “RGB color”)

- packed order
- Supports TIFF, BMP, JPEG, PNG and RAW

Indexed images (“8-bit color”)

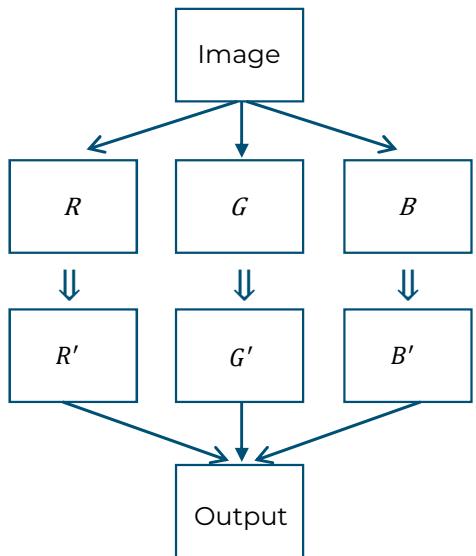
- Up to 256 colors max (8 bits)
- Supports GIF, PNG, BMP and TIFF (uncompressed)



General Strategies for Processing Color Images

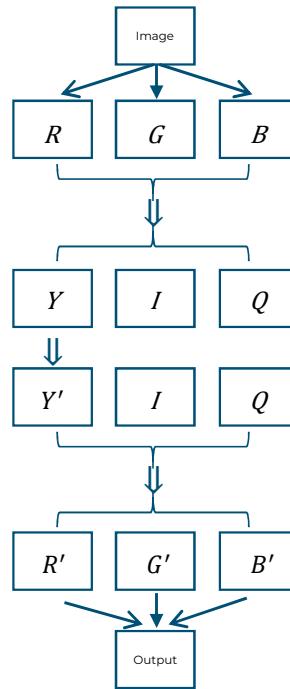
Strategy 1:

Process each RGB matrix separately



Strategy 2:

Compute luminance (weighted average of RGB), process intensity matrix



RGB to HSV Conversion

- Define the following values

$$C_{\text{high}} = \max(R, G, B), C_{\text{low}} = \min(R, G, B), C_{\text{rng}} = C_{\text{high}} - C_{\text{low}}$$

- Find saturation of RGB color components ($C_{\text{max}} = 255$)

$$S_{\text{HSV}} = \begin{cases} \frac{C_{\text{high}}}{C_{\text{max}}} & C_{\text{high}} > 0 \\ 0 & \text{otherwise} \end{cases}$$

- Find luminance value

$$V_{\text{HSV}} = \frac{C_{\text{high}}}{C_{\text{max}}}$$

RGB to HSV Conversion

- Normalize each component using

$$R' = \frac{C_{high} - R}{C_{rng}} \quad G' = \frac{C_{high} - G}{C_{rng}} \quad B' = \frac{C_{high} - B}{C_{rng}}$$

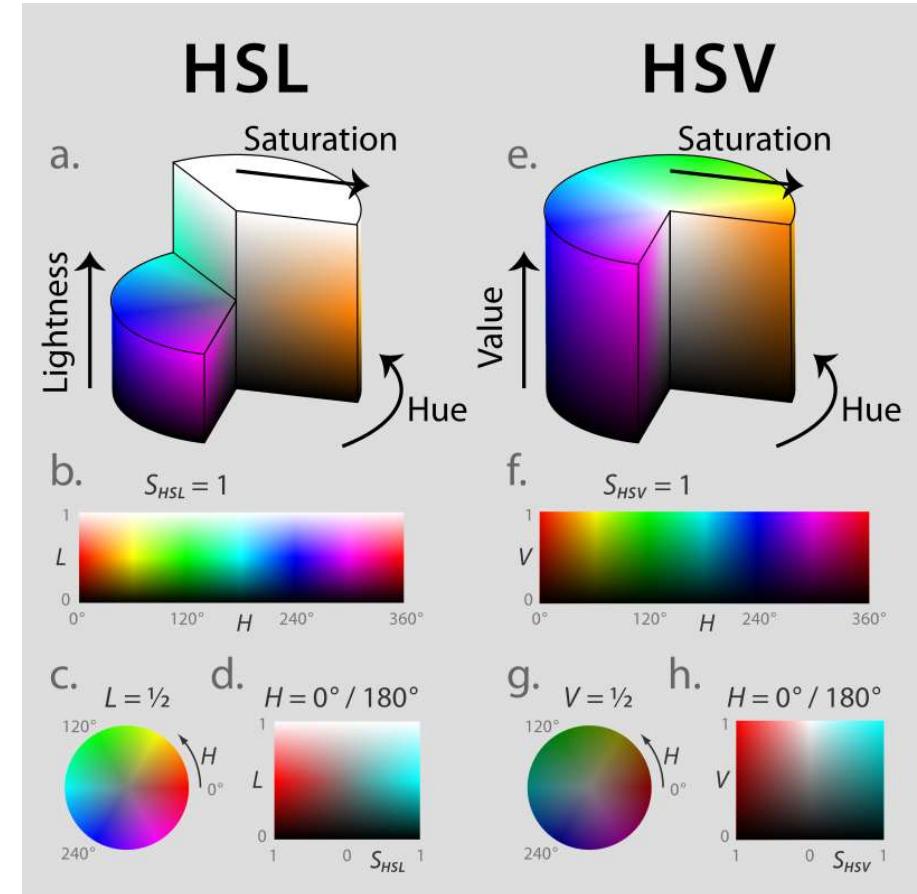
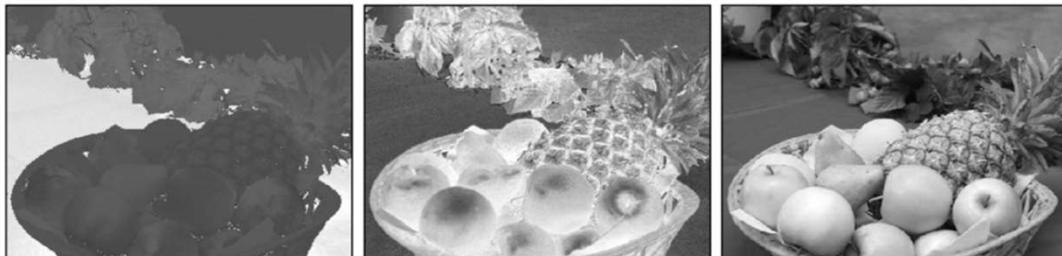
- Calculate preliminary hue value H' as

$$H' = \begin{cases} B' - G' & \text{if } R = C_{high} \\ R' - B' + 2 & \text{if } G = C_{high} \\ G' - R' + 4 & \text{if } B = C_{high} \end{cases}$$

- Finally, obtain final **hue** value by normalizing to interval [0,1]

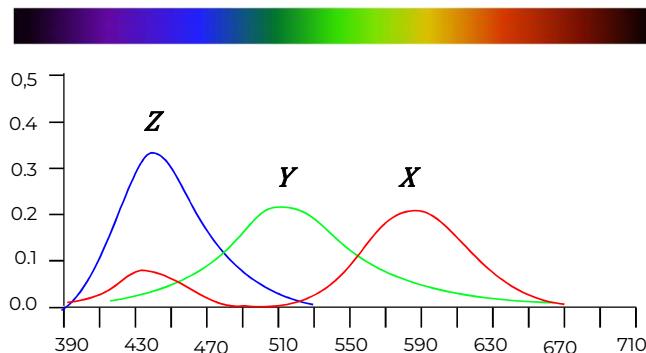
$$H_{HSV} = \frac{1}{6} \cdot \begin{cases} (H' + 6) & \text{for } H' < 0 \\ H' & \text{otherwise} \end{cases}$$

RGB to HSV (Hue, Saturation, Value)



CIE Color Space

- *CIE* (Commission Internationale d'Eclairage) came up with 3 hypothetical lights X, Y, and Z with these spectra:

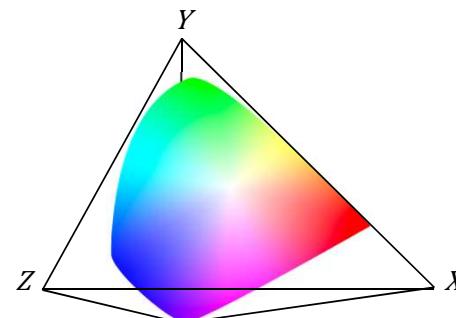


Note that

X ~ R

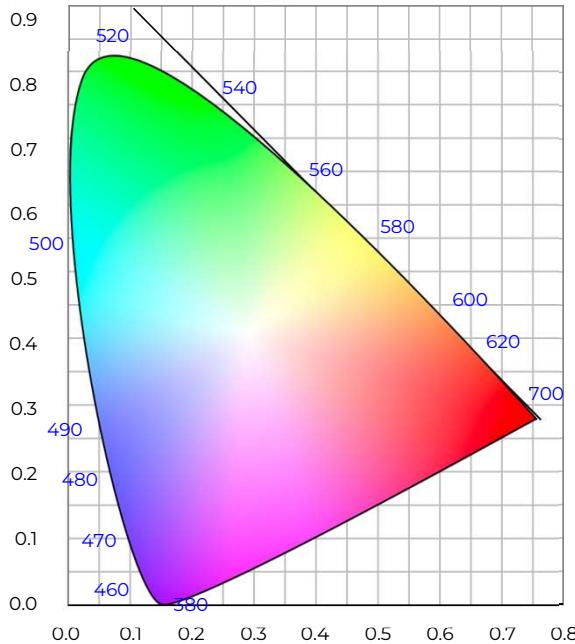
Y ~ G

Z ~ B



- **Idea:** any wavelength λ can be matched perceptually by *positive* combinations of X,Y,Z
- CIE created table of XYZ values for all *visible* colors

CIE Chromaticity Diagram (1931)



- For simplicity, we often project to the 2D plane
- Also normalize

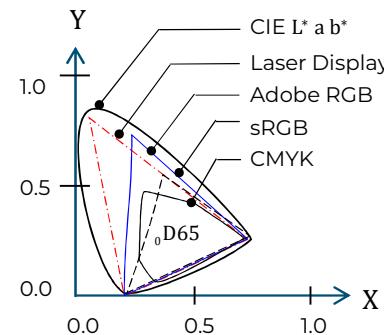
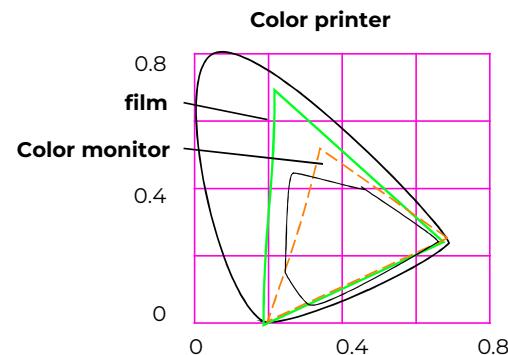
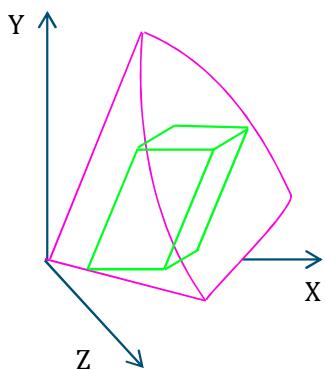
$$x + y + z = 1$$

$$x = \frac{X}{X + Y + Z} \quad y = \frac{Y}{X + Y + Z} \quad z = \frac{Z}{X + Y + Z}$$

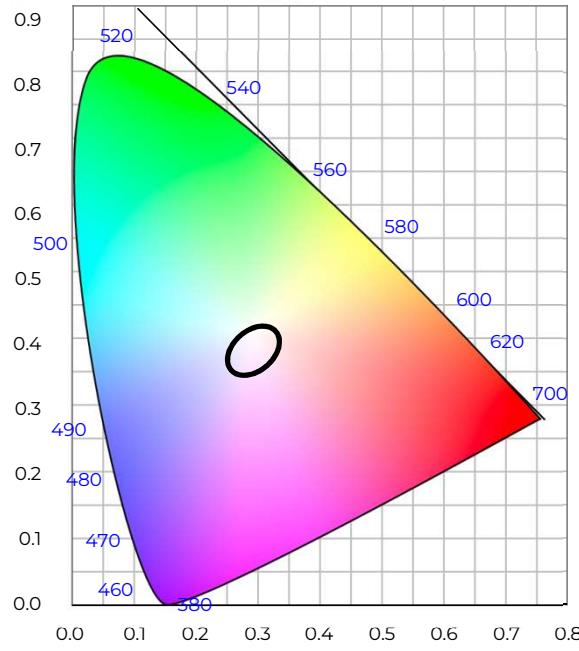
- **Note:** Inside horseshoe visible, outside invisible to eye
- **Note:** Look up x, y
Calculate z as $1 - x - y$

Device Color Gamuts

- The RGB color cube sits within CIE color space
- We can use the CIE chromaticity diagram to compare the gamuts of various devices
- E.g. compare color printer and monitor color gamuts



CIE Color classification



Human Vision is Trichromatic

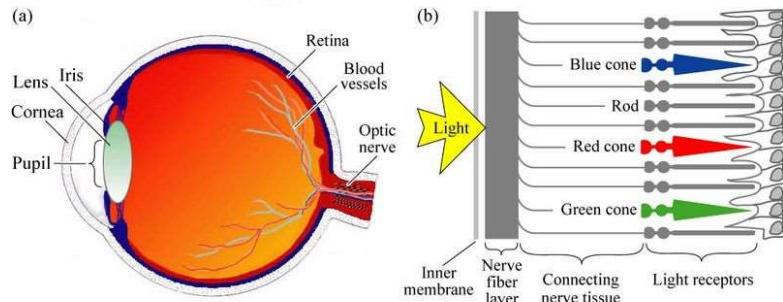
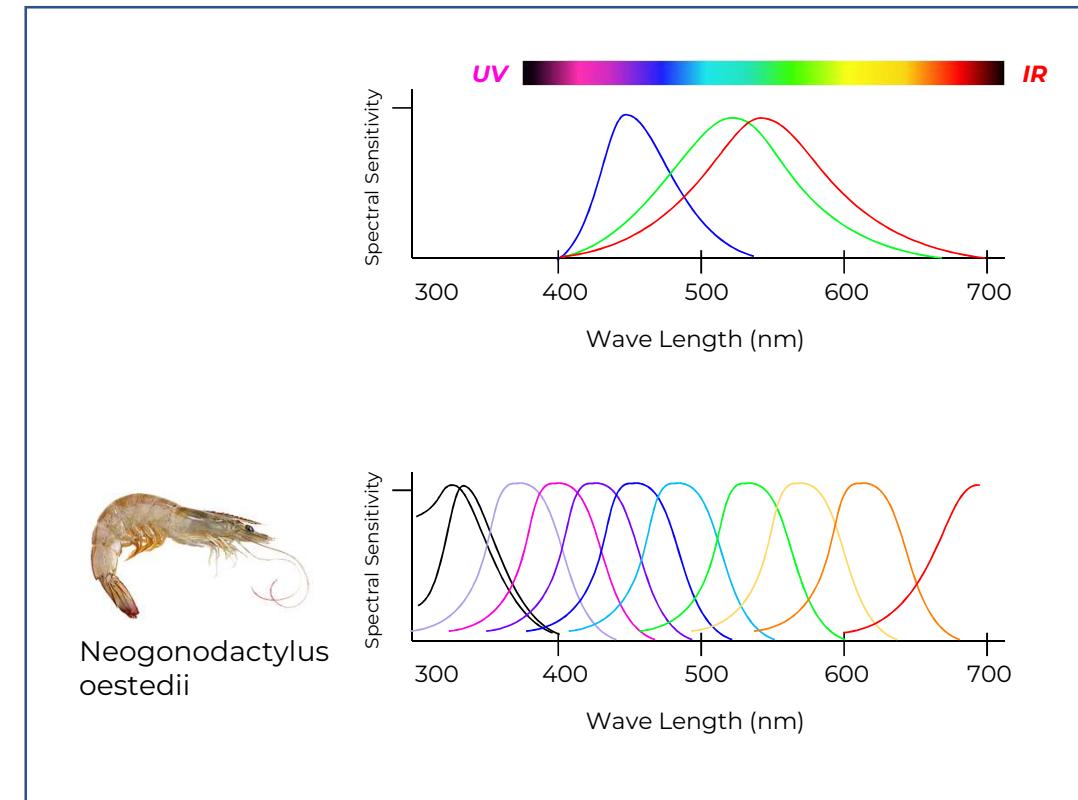


Fig. 16.1. (a) Cross section through a human eye. (b) Schematic view of the retina including rod and cone light receptors (adapted from Encyclopedia Britannica, 1994).

E. F. Schubert
Light-Emitting Diodes (Cambridge Univ. Press)
www.LightEmittingDiodes.org

Shrimp is better!



Practical Machine Vision HOW TO's



- **Camera selection:**

<https://www.baslerweb.com/en/vision-campus/vision-systems-and-components/camera-selection/>

- **Lens & Lighting Selection:**

<https://www.baslerweb.com/en/vision-campus/vision-systems-and-components/find-the-right-lens/>

Resources

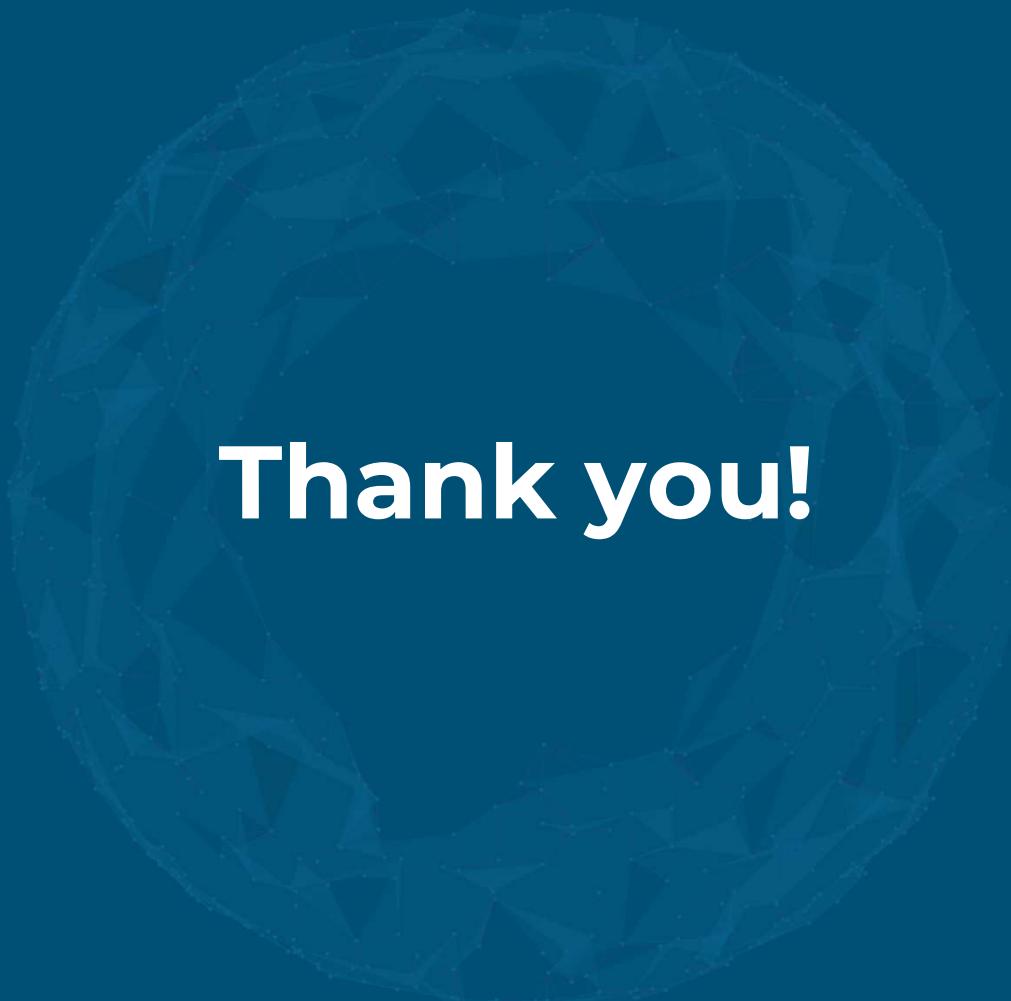
- Richard Szeliski, "Computer Vision" ISBN-13: 978-1848829343
- Rastislav Lukac, "Computtional Photography" ISBN-13: 978-1439817490
- Rick Nolthenius, Noah Snavely, George Bebis, S. Seitz, L. Zhang, D. Lowe, L. Fei-Fei, Antony Lam
- Pics: Wiki commons & Khan academy
- <http://cvcl.mit.edu/hybridimage.htm>
- MIT media lab materials, MERL edge detection
- https://vas3k.ru/blog/computational_photography/
- <https://www.baslerweb.com/en/vision-campus/vision-systems-and-components/find-the-right-lens/>
- <http://www.pptsing.com>
- <https://www.cambridgeincolour.com/tutorials/cameras-vs-human-eye.htm>
- <https://ai.googleblog.com/2014/10/hdr-low-light-and-high-dynamic-range.html>
- <http://web.media.mit.edu/~raskar/Talks/ETCVparis08/raskarCompPhotoEpsilonCodedETVC08paper.pdf>
- <https://graphics.stanford.edu/talks/compphot-publictalk-may08.pdf>
- <https://www.cambridgeincolour.com/tutorials/cameras-vs-human-eye.htm>

© Credits

- **Slide |2, 38, 47|** Image credit by MetalSunde94;
- **Slide |3,4,6|** [https://en.wikipedia.org/wiki/The_Tower_of_Babel_\(Bruegel\)](https://en.wikipedia.org/wiki/The_Tower_of_Babel_(Bruegel)) Skolkovo Institute of Science and Technology;
- **Slide |9, 10|** Albert Einstein, official 1921 Nobel Prize in Physics photograph
- **Slide |14|** <http://www.normankoren.com/Tutorials/MTF.html> photo by Andrew Hutchings
- **Slide |20|** <https://cheezburger.com/6847687168> photo by Denise Johnson
- **Slide |30-34|** https://en.wikipedia.org/wiki/Saint_Basil%27s_Cathedral#/media/File:St._Basil_Cathedral.jpg photo by Anton Zelenov
- **Slide |54|** photo by Craig McKay
- **Slide |56-59|** photo by Jordan Whitfield
- **Slide |60,65|** Albert Einstein, official 1921 Nobel Prize in Physics photograph
- **Slide |63|** https://en.wikipedia.org/wiki/Palm_cockatoo
- **Slide |66|** Lena Söderberg
- **Slide |67|** source lazebnik
- **Slide |68|** https://en.wikipedia.org/wiki/Pattern#/media/File:Aloe_polyphylla_spiral.jpg
- **Slide |71|** Boston City Hall, designed by Kallmann, McKinnell, & Knowles
- **Slide |77|** Credit W. Freeman et al, "Computer Vision for Interactive Computer Graphics", IEEE Computer Graphics and Applications, 1998
- **Slide |80|** Fergus, et al. "Removing Camera Shake from a Single Photograph", SIGGRAPH 2006, MIT Ray and Maria Stata Center

© Credits

- **Slide |3,23|** Image credit by Stevacer
- **Slide |14|** <https://www.wikihow.com/Use-Your-Digital-Camera%27s-ISO-Setting>
- **Slide |22|** <https://www.pyimagesearch.com/2018/12/17/image-stitching-with-opencv-and-python>
- **Slide |24|** kainz <https://ai.googleblog.com/2014/10/hdr-low-light-and-high-dynamic-range.html>
- **Slide |26|** ai.googleblog.com
- **Slide |27|** Ignatov, et al, ICCV2017: arXiv:1704.02470, 2017
- **Slide |29|** Yuan SIGGRAPH 2007
- **Slide |32, 49|** Lena Söderberg
- **Slide |33|** Good collection of kernels in python: <https://gist.github.com/jdumont0201/4327a1c26293b44e8d9c1cf1deef25bb>
- **Slide |47|** Source: Wikipedia
- **Slide |55|** Source: S. Seitz
- **Slide |58|** <https://www.buzzfeed.com/>
- **Slide |60|** https://vas3k.ru/blog/computational_photography/
- **Slide |67|** source lazebnik
- **Slide |76|** E F Schubert Light Emitting Diodes (Cambridge Univ Press <https://www.freepngs.com/shrimp-pngs?pgid=ivyg5wny-25516ff4-5461-11e8-a9ff-063f49e9a7e4>
- **Slide |91|** Ramesh
- **Slide |102|** Albert Einstein, official 1921 Nobel Prize in Physics photograph



Thank you!

TV Color Spaces – YUV, YIQ, YC_bC_r

- YUV, YIQ: color encoding for analog NTSC and PAL
- YC_bC_r : Digital TV encoding
- Key common ideas:
 - Separate luminance component Y from 2 chroma components
 - Instead of encoding colors, encode color differences between components (maintains compatibility with black and white TV)

YUV color space

- Basis for color encoding in analog TV in North America (NTSC) and Europe (PAL)
- Y components computed from RGB components as

$$Y = 0.299 \cdot R + 0.587 \cdot G + 0.114 \cdot B$$

- UV components computed as

$$U = 0.492 \cdot (B - Y) \text{ and } V = 0.877 \cdot (R - Y)$$

YUV color space

- Entire transformation from RGB to YUV

$$\begin{pmatrix} Y \\ U \\ V \end{pmatrix} = \begin{pmatrix} 0.299 & 0.587 & 0.114 \\ -0.147 & -0.289 & 0.436 \\ 0.615 & -0.515 & -0.100 \end{pmatrix} \cdot \begin{pmatrix} R \\ G \\ B \end{pmatrix}$$

- Invert matrix above to transform from YUV back to RGB

$$\begin{pmatrix} R \\ G \\ B \end{pmatrix} = \begin{pmatrix} 1.000 & 0.000 & 0.140 \\ 1.000 & -0.395 & -0.581 \\ 1.000 & -0.515 & -0.100 \end{pmatrix} \cdot \begin{pmatrix} Y \\ U \\ V \end{pmatrix}$$